

Closing the Airbnb Pricing Adoption Gap

An Interpretable BI + DS Decision-Support Tool

Helping Airbnb's host-success and pricing teams identify which listings to nudge, by how much, and why — with model-backed explanations hosts will actually trust.

For: Airbnb Pricing Team | Host-Success Team | Marketplace Leadership

Team 3 | Ghilliann Lou · Jianxiong Zhu · Ruimin Pei · Thomas Zhang · Yuxuan Jin

Business Problem: The Pricing Adoption Gap

The Problem

Airbnb's Smart Pricing exists, but adoption is stuck in the **single-digit to low-teens percent range**.

Hosts distrust opaque automated changes they cannot interpret.

Result:

34.8% of listings are underpriced by >15%, leaving revenue on the table.

Why Now

Post-recovery pricing (2023) is structurally higher than 2019. The opportunity window to close adoption gaps is now — before competitors build trust-first recommendation layers.

34.8%

listings underpriced >15%

\$1.28M

est. annual GBV uplift at 25% adoption

\$48.99/night

avg gap per flagged listing

\$179K

est. annual commission uplift (14%)

Cost of inaction: every quarter without trust-based recommendations leaves millions in GBV and commission uncaptured.

Stakeholders · Decision · Target Variable · KPIs

Decision: Which 1,000 listings should the host-success team contact next quarter, what should each host be told, and why?

Stakeholder	Role	Dashboard View
Marketplace Strategy / Host-Success	Primary decision-maker	Top Underpriced Queue
Host-Success Reps	Operational (customer-facing)	Listing Recommendation
Pricing / Analytics Team	Internal data consumers	Model Performance Tab
Hosts (indirect)	Beneficiaries (~35% underpriced)	Recommendations

Model KPIs

- MAE (primary) – \$/night error
- RMSE – penalizes large errors
- MAPE – relative % error
- R^2 / Adjusted R^2 – variance explained

Business KPIs

- GBV uplift (adoption sensitivity 10/25/50%)
- Commission uplift at 14% rate
- Recommendation adoption rate
- Priority queue efficiency (listings / FTE)

Target Variable: Nightly listing price (continuous, modeled in log space) | Classification: Underpriced (<-15%) / Fairly Priced / Overpriced (>+15%)

Data Sources: Multi-Year, Multi-Market Coverage

Dataset	Scope	Size	Role in Project
AB_NYC_2019.csv	New York City	~48K listings	Pre-pandemic urban baseline; dense market
AB_US_2020.csv	National (US)	~12K listings	Pandemic-era snapshot; depressed price regime
AB_US_2023.csv	National (US)	~19K listings	Post-recovery national pricing signal
listings-Boston.csv	Boston metro	~3.4K listings	Rich amenity/review features; city-level depth
listings-Seattle.csv	Seattle metro	~6.8K listings	Rich amenity/review features; Pacific NW market

Why This Mix of Datasets Is Appropriate

Temporal Coverage:

Pre-pandemic (2019), pandemic (2020), and post-recovery (2023) regimes allow regime-aware modeling via `data_year` feature

Geographic Breadth:

NYC and US national files provide volume; Boston and Seattle add listing-level feature richness (amenities, reviews)

Schema Harmonization:

Different schemas unified: `neighbourhood/neighbourhood_cleansed`, city inferred from filename, missing columns filled with sentinels

Stratified Sampling:

Down-sampled to 25K rows stratified by city so smaller markets (Boston, Seattle) are not drowned out by NYC volume

Data Preparation & Feature Engineering

Data Cleaning & Harmonization

- Price parsing: strip \$ and commas; coerce to float
- Outlier cap: 99th percentile (\$1,326) — data-driven, not hardcoded
- Room type: restrict to 4 canonical types; drop unknown
- Neighbourhood harmonization: neighbourhood_cleansed → neighbourhood
- City inference: from filename when column absent
- Missingness diagnostic: missing review_scores_rating → \$+9 price delta; treated as signal via 'Unknown' category
- Stratified down-sample: 25K rows stratified by city

Engineered Features

availability_ratio:

availability_365 / 365 — booking pressure signal

review_activity:

reviews_per_month — listing popularity proxy

amenity_indicators:

Binary flags: wifi, kitchen, AC, parking, pool, gym, pets

amenity_score:

Weighted sum — pool×2, gym×1.5, others×1

has_luxury_amenities:

Binary: pool OR gym OR parking

host_scale_bucket:

Bucketized host listings count (single/small/medium/large)

city_roomtype_segment:

City × room_type interaction — local pricing pattern

data_year:

Regime-shift awareness: 2019/2020/2023

EDA: The Business Insights That Drive the Pricing Tool

EDA Q2: How does median price vary across the top 12 markets?



📍 Median price spans \$69–\$150 (2.2×) across 12 markets — location is the #1 pricing driver.

Takeaway: Median price differs materially across years. By including 'data_year' as

EDA Q5: How does price vary by room type?

```
/tmp/ipykernel_1852/585967111.py:146: MatplotlibDeprecationWarning: The 'labels' pa  
bp = ax.boxplot()
```



🏠 Entire homes earn 2–3× the median of private rooms — room type is the strongest categorical predictor.

🛎️ Amenity premiums vary by city

Pool adds +\$47–\$49 overall, but premium is market-specific — some amenities show negative correlation in lower-priced markets. This argues for model-learned, not rule-based, adjustments.

⚠️ 34.8% of listings systematically mispriced

At 15% threshold, 1,741 underpriced listings flagged with avg gap of \$48.99/night. At 5–10%, we mostly catch modeling noise. 15% is the signal-to-noise sweet spot.

Modeling Approach: Why It Is Credible

- | | | |
|---|---|---|
| 1 | Naive Baseline
Segment median (city × room_type × bedrooms) | <i>Floor — any ML model must beat this</i> |
| 2 | Linear Regression
Log-price OLS with preprocessed features | <i>M3 baseline tier — interpretable benchmark</i> |
| 3 | Random Forest
300 trees, max_depth 18 (bagging) | <i>Reduces variance; handles non-linearities</i> |
| 4 | Gradient Boosting
200 estimators, lr 0.05 (sklearn) | <i>Reduces bias; sequential improvement</i> |
| 5 | XGBoost ★
500 trees, depth 7, lr 0.03 — GridSearchCV tuned | <i>M3 recommended; best validation MAE</i> |

Validation: 60/20/20 stratified train/val/test split + 3-fold cross-validation | GridSearchCV with neg_mean_absolute_error | Final metrics reported once on held-out test set

Model Results, Interpretability & Key Findings

Test-Set Model Comparison

Model	MAE	R ²	MAPE
Naive Baseline	\$71.19	0.09	41.3%
Linear Regression	\$68.95	-2.02	42.9%
Gradient Boosting	\$61.51	0.316	36.9%
Random Forest	\$58.57	0.384	35.3%
XGBoost (Tuned) ★	\$58.15	0.386	34.5%


\$58.15

MAE (test set)

18.3%

Better than naive baseline

Top Permutation Importances (model-agnostic)



Permutation interpretation: shuffling a feature's values across rows breaks its relationship with the target variable.

- Room type dominates: shuffling room_type raises MAE by \$22/night — the single strongest predictor
- Location pair matters: lat/lon together add \$14/night to error — sub-mile pricing clusters confirmed
- Adj R² = 0.324 — honest about ~67% unexplained variance (photography, host responsiveness, calendar)

Dashboard: How a Host-Success Rep Uses the Tool

1

Filter Market

Select city & room type in sidebar

3

Click a Listing

Tab 1 auto-populates with inputs

5

See SHAP Drivers

'Why This Price?' panel — top 3 reasons to use on calls

2

View Priority Queue

Tab 2: ranked underpriced listings by estimated uplift

4

Read Recommendation

Banner: Underpriced / Overpriced / Fairly Priced + % gap

6

Export & Track

Download top-N CSV for outreach tracking
(top_500_pilot_listings.csv)

Key Dashboard Outputs per Listing:

- Predicted market price
- Actual listed price
- Mispricing % + label
- Estimated annual uplift
- SHAP top-3 explanation
- Threshold slider (10–25% by segment)


Business Impact: Adoption-Rate Sensitivity

GBV & Commission Uplift from Underpriced Listings (5,000-listing test set)


Adoption Rate	Adjusted GBV Uplift	Commission Uplift (14%)	Listings Flagged
10%	\$511,778	\$71,649	1,741
25% (central)	\$1,279,446	\$179,122	1,741
50%	\$2,558,892	\$358,245	1,741

Held constant: 60 booked nights/yr · 14% Airbnb commission · Extrapolates to full marketplace footprint at materially larger scale

Top Cities by Commission Uplift (25% Adoption)



Threshold Sensitivity (coverage vs signal)



Defence of 15%: at 5% and 10%, the average \$ gap per flagged listing is small enough

Recommendations, Expected Impact & Next Steps

01

Launch 90-Day A/B Pilot (Manhattan)

Population: NYC entire-home underpriced $\geq 15\%$, top ~ 500 listings
Intervention: host-success rep contacts with SHAP-derived top-3 reasons
Decision rule to scale: treatment GBV/listing > control by 3% AND adoption > 15%

02

Build Explanation Layer into Host-Facing Dashboard

SHAP-based 'Why This Price?' panel makes every recommendation defensible
Location, room type, amenity contributions in plain language
Highest-leverage product lever for adoption — this is product work, not modeling

03

Operate with Segment-Specific Thresholds

High-confidence markets (Boston $\sim \$42$ MAE, Seattle $\sim \$27$ MAE): use 10% threshold
High-error markets (Austin $\sim \$79$ MAE, SF $\sim \$68$ MAE): use 20–25% threshold
Threshold slider in dashboard — ops manager sets per segment



Limitations

Listed price $\neq$ realized booking price · Observational (not causal) data · 2019–2023 temporal regime · Sparse-segment uncertainty · Fairness audit required before rollout



Future Roadmap

Seasonality-aware dynamic pricing · Booking-rate (realized revenue) model · Host tenure / fairness analysis · City-specific model deployment · Production drift monitoring